

Operating system 10 marks answer key

MAY 2021

PART C — (3 × 10 = 30 marks)

Answer any THREE questions.

20. Explain briefly different types of OS.
21. What is virtual memory? Explain.
22. Discuss about deadlock avoidance in detail.
23. Briefly Explain free space management.
24. Explain different file operations in detail.

Answer

20. **Different Types of Operating Systems:**

There are various types of operating systems designed for different purposes. Here are some of the most common types:

- **Real-Time Operating Systems (RTOS):** These are designed for systems where the timely execution of tasks is critical, such as in embedded systems, aerospace, and medical devices. RTOS guarantees tasks are executed within specified time constraints.
- **Single-User, Single-Tasking OS:** These are simple OSes that run on embedded systems or very basic computers, where only one user can perform one task at a time.
- **Single-User, Multi-Tasking OS:** These OSes allow a single user to perform multiple tasks concurrently. Examples include older versions of Microsoft Windows (e.g., Windows 3.1).
- **Multi-User OS:** These OSes allow multiple users to interact with the system simultaneously, each having their processes. Examples include various versions of Unix and Linux.
- **Multi-Tasking OS:** Multi-tasking OSes allow multiple processes or tasks to run concurrently. Modern desktop and server OSes like Windows, macOS, and Linux fall into this category.
- **Distributed OS:** Distributed OSes are designed for systems where multiple computers are networked and function as a single system. They are used in cloud computing and large-scale data centers.

- **Network OS:** These OSes primarily facilitate network-related functions, managing network connections and data transfers.

- **Mobile OS:** Operating systems for mobile devices such as Android, iOS, and Windows Mobile.

- **Embedded OS:** These are lightweight, specialized OSes designed for specific embedded systems, such as IoT devices and consumer electronics.

Each type serves specific needs and may come with unique characteristics and features tailored to their use cases.

21. **Virtual Memory:**

Virtual memory is a memory management technique used in modern operating systems. It provides an illusion of a larger memory space to applications than the physical memory (RAM) available in the system. This is achieved by using a combination of RAM and disk storage. Here's how it works:

- The OS divides the logical memory (used by applications) into fixed-sized blocks called pages.
- Physical memory (RAM) is divided into blocks of the same size called frames.
- The OS maintains a page table that maps pages to frames.
- When a process references a page not currently in RAM, a page fault occurs. The OS swaps a page from RAM to disk to make room for the new page.
- Virtual memory provides numerous advantages, such as efficient use of memory, isolation, protection, and ease of memory management.

22. **Deadlock Avoidance:**

Deadlock avoidance is a strategy to prevent deadlocks from occurring in a system. It works by ensuring that at least one of the necessary conditions for deadlock cannot hold. One popular method is using the Banker's algorithm:

- Each process must declare its maximum resource requirements when it starts.
- The system keeps track of the available resources.
- When a process requests resources, the system checks if it can satisfy the request without leading to a deadlock.
- If the request is safe, the resources are allocated; otherwise, the process must wait.

The Banker's algorithm prevents deadlock by ensuring that, even with all processes requesting resources at the same time, there is a safe sequence of resource allocation. If no safe sequence is possible, the request is denied.

23. **Free Space Management:**

Free space management is the process of managing and tracking available space on storage devices, primarily in file systems. There are several methods for managing free space:

- **Bit Vector:** Each block is represented by a bit in a bit vector, where 1 means the block is allocated, and 0 means it's free.
- **Linked List:** Blocks are linked in a linked list, with each block pointing to the next free block.
- **Grouping:** Blocks are grouped into clusters or extents. A cluster is marked as free when all blocks within it are free.
- **Bitmap:** Similar to a bit vector, but uses bitmap data structures for efficiency.

Efficient free space management is crucial for file systems to allocate and deallocate storage space effectively and minimize fragmentation.

24. **Different File Operations:**

File operations are fundamental to file systems and include:

- **Create:** This operation involves creating a new file in the file system. The operating system allocates space for the file and sets up metadata.
- **Open:** Opening a file allows processes to access its content. The OS checks permissions and manages file pointers.
- **Read:** Reading from a file retrieves data from it and transfers it to the requesting process.
- **Write:** Writing to a file allows processes to modify its content. The OS handles data storage and ensures changes are persistent.
- **Close:** Closing a file releases resources and any locks, ensuring the file is no longer in use by the process.
- **Delete:** Deleting a file removes it from the file system, freeing up the space it occupied.

These operations are critical for managing data in a file system, and the OS must ensure that they are performed safely and efficiently.

MAY 2022

PART C — (3 × 10 = 30 marks)

Answer any THREE questions.

20. Discuss on the different types of operating systems.
21. Explain any four types of CPU scheduling algorithms.

22. Describe the demand paging in memory management techniques.
23. Explain in detail about disk space allocation methods.
24. Explain the distinct feature of I/O systems.

Answers

20. **Different Types of Operating Systems:**

Operating systems come in various types, each designed for specific use cases. Let's explore a few of them:

- **Real-Time Operating Systems (RTOS):** These OSes are optimized for systems where precise timing and responsiveness are crucial. They are commonly used in embedded systems, automotive control systems, medical devices, and industrial automation. RTOS ensures that tasks are executed within strict time constraints, making them ideal for applications where missed deadlines can have serious consequences.

- **Single-User, Single-Tasking OS:** This type of OS is found in extremely basic or legacy systems, where only one user can perform a single task at a time. It's not suitable for modern multi-tasking environments and is rarely used today.

- **Single-User, Multi-Tasking OS:** These OSes allow a single user to run multiple applications concurrently. An example is early versions of Microsoft Windows, like Windows 3.1.

- **Multi-User OS:** Designed for systems with multiple users, multi-user OSes allow multiple users to interact with the system simultaneously. They provide features for user account management, file sharing, and security. Variants of Unix, Linux, and modern Windows are typical examples.

- **Multi-Tasking OS:** These are prevalent in most modern computers and servers. Multi-tasking OSes enable multiple processes or applications to run concurrently. They employ scheduling algorithms to allocate CPU time fairly and efficiently to running processes.

- **Distributed OS:** Distributed OSes manage a network of interconnected computers as a single system, allowing processes to communicate and share resources across the network. They are used in data centers and cloud computing environments.

- **Network OS:** Network OSes focus on managing network-related tasks, such as network protocol handling and data routing. Cisco's IOS (Internetwork Operating System) used in network routers is an example.

- **Mobile OS:** These operating systems are specifically designed for mobile devices, such as Android and iOS for smartphones and tablets. They incorporate features like touchscreen support, power efficiency, and app stores.

- **Embedded OS:** Embedded systems, like IoT devices, consumer electronics, and industrial machines, often use lightweight embedded OSes tailored to specific hardware and functionality. Examples include FreeRTOS and VxWorks.

Each type of operating system serves distinct purposes and is tailored to meet the specific requirements of the target application or hardware.

21. **CPU Scheduling Algorithms:**

CPU scheduling algorithms are used by the operating system to determine which process to execute next. Four common scheduling algorithms include:

- **First-Come, First-Served (FCFS):** This is a non-preemptive scheduling algorithm where the process that arrives first is assigned the CPU first. It's simple but may lead to the "convoy effect" where a long process blocks shorter ones behind it.

- **Shortest Job First (SJF):** SJF is a non-preemptive or preemptive scheduling algorithm where the process with the shortest burst time is selected next. It minimizes the average waiting time and turnaround time but requires knowledge of process execution times, which is often not available.

- **Round Robin (RR):** Round Robin is a preemptive scheduling algorithm that assigns each process a fixed time quantum, and they take turns executing. It's fair and prevents long processes from blocking others, but it can lead to overhead due to context switches.

- **Priority Scheduling:** In priority scheduling, each process is assigned a priority, and the process with the highest priority is executed first. Preemptive versions can suffer from priority inversion issues. It's essential to avoid starvation and ensure proper priority assignment.

These algorithms vary in terms of fairness, efficiency, and suitability for different scenarios. Choosing the right algorithm depends on the system's requirements and goals.

22. **Demand Paging in Memory Management:**

Demand paging is a memory management technique used by modern operating systems. It combines the concepts of virtual memory and paging. Here's how it works:

- A process's entire address space is initially stored on secondary storage (typically a hard drive or SSD), not in physical RAM.

- When a process is loaded into memory, only the necessary pages are brought in from secondary storage, rather than loading the entire process.

- As a process runs, additional pages are fetched from secondary storage as needed. Pages that are not accessed remain on disk.

- Demand paging allows more processes to run concurrently since only the active parts of each process are loaded into RAM.

- It leverages the benefits of virtual memory, including process isolation and protection.

Demand paging increases memory utilization and can improve system performance by reducing the amount of physical memory required to run multiple processes simultaneously.

23. **Disk Space Allocation Methods:**

Disk space allocation methods are used to manage the allocation and organization of data on storage devices, such as hard drives. Several methods exist:

- **Contiguous Allocation:** In this method, files are stored in contiguous blocks on the disk. It's straightforward and efficient for reading, but it leads to fragmentation and can be problematic when files need to grow.
- **Linked Allocation:** Each file is a linked list of blocks. This method eliminates fragmentation but can lead to inefficient disk I/O due to scattered data blocks.
- **Indexed Allocation:** A separate index block holds pointers to data blocks. It efficiently supports direct access to data, reducing I/O overhead. It's commonly used in file systems like NTFS.
- **Multi-Level Indexing:** This method extends indexed allocation with multiple levels of indexing to support large files efficiently.
- **Bitmap Allocation:** In this approach, a bitmap is used to manage the allocation status of each block. It's efficient for both allocation and deallocation and is used in modern file systems like FAT and NTFS.

Each allocation method has its advantages and disadvantages, and the choice depends on the requirements of the file system and the use case.

24. **Distinct Features of I/O Systems:**

Input/Output (I/O) systems play a crucial role in enabling data exchange between an operating system and external devices. They possess distinct features to ensure efficient data transfer:

- **Abstraction:** I/O systems abstract the low-level details of device communication, allowing applications to use a standardized interface for all types of devices.
- **Buffering:** Buffering involves temporarily storing data in memory before sending it to a device or after reading it from a device. It helps smooth data transfer and reduces latency.
- **Caching:** Caching involves keeping frequently accessed data in a cache memory to reduce the need for repeated disk or device accesses. Caching improves I/O performance significantly.
- **Direct Memory Access (DMA):** DMA allows devices to access system memory directly without involving the CPU. This offloads processing overhead from the CPU, improving data transfer rates.

- **Interrupts:** I/O systems use interrupts to notify the CPU when a device operation is complete, enabling asynchronous processing and efficient multitasking.
- **Error Handling:** Robust error handling mechanisms are in place to detect and manage errors during data transfer, ensuring data integrity.
- **File System Interaction:** I/O systems facilitate interaction with the file system, allowing data to be read from and written to files.

I/O systems are critical to the overall performance and functionality of an operating system, as they enable communication with a wide range of devices and external storage.

These features collectively ensure efficient, reliable, and high-performance data transfer between the operating system and external devices, which is essential for a well-functioning computer system. The various I/O management techniques and strategies help minimize the impact of slow I/O operations on overall system performance, ensuring a smooth user experience.

In summary, operating systems can be classified into various types, each tailored to specific use cases. CPU scheduling algorithms, such as FCFS, SJF, RR, and priority scheduling, determine the order in which processes are executed, impacting system efficiency and fairness. Demand paging is a memory management technique that optimizes RAM usage by loading only the necessary parts of a process into memory. Disk space allocation methods, including contiguous, linked, indexed, multi-level indexing, and bitmap allocation, determine how data is organized on storage devices. I/O systems exhibit unique features, such as abstraction, buffering, caching, DMA, interrupts, error handling, and file system interaction, to facilitate efficient data transfer between the operating system and external devices.

These components and techniques collectively contribute to the smooth and efficient operation of modern computer systems and are vital for managing the complex interactions between hardware and software in a way that ensures user productivity and system reliability.

DECEMBER 2022

PART C — (3 × 10 = 30 marks)

Answer any THREE questions.

20. Explain about distributed systems.
21. Explain the various CPU scheduling algorithms.
22. Explain in detail about paging.
23. Explain in detail about file allocation methods.
24. Explain about memory management.

Answer

20. **Distributed Systems:**

A distributed system is a network of autonomous computers that work together as a unified whole. In a distributed system, each computer, referred to as a node or host, has its processing power and memory. These nodes communicate and coordinate with each other to achieve a common goal. Here are some key aspects of distributed systems:

- **Autonomy:** Each node in a distributed system operates autonomously and can make local decisions. It doesn't rely on a central authority for every action.
- **Communication:** Nodes in a distributed system communicate with each other, typically using a network. This communication allows them to share information and collaborate on tasks.
- **Transparency:** A well-designed distributed system aims to hide the complexities of distribution from users and applications. Transparency includes location transparency (users don't need to know where data is stored) and access transparency (users access data in the same way, regardless of its location).
- **Fault Tolerance:** Distributed systems often incorporate redundancy and error-handling mechanisms to ensure continued operation even in the presence of hardware failures or network issues.
- **Scalability:** Distributed systems can scale horizontally by adding more nodes to handle increased workloads. This makes them suitable for large-scale applications and services.
- **Examples:** Common examples of distributed systems include the internet, cloud computing platforms, distributed databases, and peer-to-peer networks.

Distributed systems are essential for enabling modern computing, allowing applications to harness the power of multiple computers and providing resilience and scalability.

21. **Various CPU Scheduling Algorithms:**

CPU scheduling algorithms are used by operating systems to determine the order in which processes are executed on the CPU. Here are some common CPU scheduling algorithms:

- **First-Come, First-Served (FCFS):** FCFS is a non-preemptive algorithm where the process that arrives first is assigned the CPU first. It's simple but may lead to the "convoy effect" where a long process blocks shorter ones behind it.
- **Shortest Job First (SJF):** SJF is a non-preemptive or preemptive algorithm where the process with the shortest burst time is selected next. It minimizes average waiting time and turnaround time but requires knowledge of process execution times, which is often not available.

- **Round Robin (RR):** Round Robin is a preemptive scheduling algorithm that assigns each process a fixed time quantum, and they take turns executing. It's fair and prevents long processes from blocking others but can lead to overhead due to context switches.

- **Priority Scheduling:** Priority scheduling assigns a priority to each process, and the process with the highest priority is executed first. It can be either preemptive or non-preemptive. Care is needed to prevent priority inversion issues.

- **Multilevel Queue Scheduling:** In this approach, processes are divided into multiple queues with different priorities. Each queue may use a different scheduling algorithm. For example, interactive processes may be in a high-priority queue, while batch processes are in a low-priority queue.

- **Multilevel Feedback Queue Scheduling:** Similar to multilevel queue scheduling, but with the ability to move processes between queues based on their behavior. This allows for dynamic adjustment of priorities.

The choice of scheduling algorithm depends on the system's requirements and goals, balancing factors like fairness, efficiency, and response time.

22. **Paging:**

Paging is a memory management scheme used in modern operating systems. It divides physical memory (RAM) and logical memory (used by processes) into fixed-sized blocks called frames and pages, respectively. Here's how it works:

- The logical memory is divided into pages, each of the same size as frames.
- The operating system maintains a page table that maps pages to frames. This mapping enables efficient memory allocation and management.
- When a process references a page that is not in RAM, a page fault occurs. The OS swaps a page from RAM to disk to make room for the new page.
- Paging provides several advantages, such as efficient use of memory, isolation, protection, and ease of memory management.

Paging simplifies memory allocation, eliminates external fragmentation, and ensures efficient usage of physical memory. It is a key component of virtual memory systems.

23. **File Allocation Methods:**

File allocation methods are used to organize and allocate space for files on storage devices, such as hard drives. Some common file allocation methods include:

- **Contiguous Allocation:** In this method, each file is stored as a single, contiguous block on the disk. It's simple but can lead to fragmentation, making it challenging to allocate space for larger files.
- **Linked Allocation:** In linked allocation, each file is a linked list of blocks. This method eliminates external fragmentation but can result in inefficient disk I/O due to scattered data blocks.

- **Indexed Allocation:** Indexed allocation uses an index block to store pointers to data blocks. It provides efficient support for direct access to data blocks and is commonly used in modern file systems like NTFS.

- **Multi-Level Indexing:** Multi-level indexing extends indexed allocation by introducing multiple levels of index blocks, enabling efficient support for large files.

- **Bitmap Allocation:** Bitmap allocation uses a bitmap data structure to manage the allocation status of each block. It is efficient for both allocation and deallocation and is used in modern file systems like FAT and NTFS.

Each file allocation method has its advantages and limitations, and the choice of method depends on the specific requirements and constraints of the file system.

24. **Memory Management:**

Memory management is a fundamental component of operating systems and plays a crucial role in ensuring efficient utilization of physical memory (RAM). Here are key aspects of memory management:

- **Address Translation:** The operating system uses techniques like paging and segmentation to map logical addresses used by processes to physical addresses in RAM. This mapping enables efficient memory allocation and protection.

- **Virtual Memory:** Virtual memory allows processes to access more memory than physically available by using disk space as an extension of RAM. It provides benefits such as isolation, protection, and efficient memory utilization.

- **Memory Protection:** Memory protection mechanisms prevent processes from accessing memory regions that do not belong to them. This enhances system security and stability.

- **Page Replacement:** In a virtual memory system, when a process references a page that is not in RAM, the OS must select a page to evict and swap in the needed page. Algorithms like LRU (Least Recently Used) and FIFO (First-In, First-Out) are used for page replacement.

- **Memory Allocation:** The OS allocates memory to processes dynamically as they request it. It also handles deallocation and reclamation of memory when processes complete.

- **Shared Memory:** Memory management facilitates interprocess communication, allowing processes to share memory regions for data exchange.

Effective memory management is essential for overall system performance, as it directly impacts the ability

PART C — (3 × 10 = 30 marks)

Answer any THREE questions.

20. Discuss Operating system services in detail.
21. Explain about the System Call in detail.
22. Discuss about the deadlock avoidance in detail.
23. Explain the Page replacement algorithm in detail.
24. Explain about the I/O Systems in detail.

Answer

20. **Operating System Services:**

Operating systems provide a wide range of services to both users and applications. These services are essential for managing hardware resources, ensuring security, and enabling efficient communication. Let's delve into the key operating system services in detail:

- **Program Execution:** The OS loads programs into memory and manages their execution, ensuring that processes run in a controlled and secure environment.
- **I/O Operations:** The OS provides services for reading from and writing to devices such as disks, keyboards, and network interfaces, abstracting the complexities of hardware interaction.
- **File System Manipulation:** OS services allow users and programs to create, delete, read, write, and organize files and directories. File system management ensures data persistence and organization.
- **Communication Services:** Operating systems enable interprocess communication, facilitating data exchange between running processes. This includes services like message passing and shared memory.
- **Error Detection and Handling:** The OS monitors system operations for errors and exceptions, providing mechanisms for error reporting and recovery to prevent system crashes.
- **User Interface Services:** Operating systems offer user interfaces to interact with the system. This includes graphical user interfaces (GUIs), command-line interfaces (CLIs), and APIs for application development.
- **Security and Access Control:** OS services implement security measures to control access to resources, authenticate users, and protect against unauthorized access.

- **Networking Services:** The OS manages network connections, supports protocols for communication over networks, and provides networking utilities for configuring and monitoring network settings.
- **Time and Date Services:** The OS keeps track of system time and provides time-related services, including clock synchronization and scheduling tasks at specific times.
- **Printing Services:** OS services handle printing tasks, managing print queues, spooling documents, and interacting with printer devices.
- **Resource Allocation:** The OS allocates resources, such as CPU time and memory, to processes to ensure efficient and fair utilization of system resources.
- **Performance Monitoring and Tuning:** OS services include tools for monitoring system performance, resource usage, and diagnosing performance issues.
- **Backup and Recovery Services:** The OS provides backup and recovery mechanisms to protect data and system configurations against loss or failure.

Operating system services are essential for creating a robust and efficient computing environment. They abstract hardware complexity, provide a consistent and secure environment, and enable user and application interaction with the system.

21. **System Call:**

A system call, often abbreviated as syscall, is a fundamental interface between applications and the operating system. System calls enable programs to request services from the OS, such as I/O operations, process control, and resource management. Here's a detailed explanation of system calls:

- **Functionality:** System calls expose OS functionality to user programs. They act as a bridge between the user space (where applications run) and the kernel space (where the operating system resides).
- **Invocation:** To make a system call, a program uses a predefined interface provided by the programming language or system libraries. In most cases, system calls are made using functions or procedures specific to the operating system.
- **Examples:** Common system calls include ``open`` and ``read`` for file operations, ``fork`` and ``exec`` for process creation, ``socket`` and ``send`` for network communication, and ``malloc`` and ``free`` for memory allocation.
- **Protection:** System calls are a controlled mechanism for interacting with the OS. They ensure that applications do not have direct access to privileged operations, which could compromise system security and stability.
- **User-Mode to Kernel-Mode Transition:** When a system call is invoked, the program transitions from user mode to kernel mode, where it gains access to protected resources.

The OS validates the request, performs the requested operation, and returns control to the user space.

- **System Call Numbers:** Each system call is associated with a unique identifier (system call number). This identifier is used to look up the corresponding kernel function for execution.

- **Error Handling:** System calls can return error codes to signal exceptional conditions. Programmers need to check these codes to handle errors gracefully.

System calls are the building blocks of system functionality, enabling applications to perform essential tasks. They are at the core of the interaction between software and the operating system, providing a controlled and secure way to harness the power of the underlying hardware and OS services.

22. **Deadlock Avoidance:**

Deadlock avoidance is a strategy used to prevent deadlock situations in a concurrent system. Deadlock occurs when processes are unable to proceed because they are waiting for resources held by other processes. Here is a detailed explanation of deadlock avoidance:

- **Resource Allocation Graph:** Deadlock avoidance techniques often use a resource allocation graph to represent the state of resource allocation in the system. In this graph, processes are nodes, resource types are represented as boxes, and resource instances are represented as dots within the boxes.

- **Resource Requests:** When a process requests a resource, it makes a request edge to the corresponding resource node in the graph. When it releases a resource, it releases the corresponding edge.

- **Deadlock Detection:** The system periodically examines the resource allocation graph to determine if any cycles (deadlock) exist. A cycle indicates that processes are waiting for resources in a way that forms a closed loop.

- **Safe and Unsafe States:** The system defines states as safe or unsafe based on the presence or absence of cycles in the resource allocation graph. Safe states have no cycles, and in such states, it is guaranteed that no deadlock can occur.

- **Resource Allocation Policy:** In a deadlock avoidance strategy, the system may implement policies to ensure that resource allocations lead to safe states. This may involve making resource allocation decisions that avoid circular wait, one of the necessary conditions for deadlock.

- **Wait-Die and Wound-Wait Schemes:** These are two common deadlock avoidance schemes for managing resource requests. Wait-Die allows older processes to wait for resources, while younger processes are aborted if they request a resource held by an older process. Wound-Wait allows younger processes to wait and older processes to be preempted if necessary.

- **Resource Allocation Algorithms:** Specific resource allocation algorithms, such as the Banker's algorithm, are used to ensure that resource allocation requests are made in a way that avoids deadlock.

Deadlock avoidance aims to ensure that the system remains in a safe state, preventing deadlock from occurring. This approach requires careful resource management and allocation policies to guarantee that resources are allocated in a manner that avoids circular wait and satisfies other necessary conditions for deadlock prevention.

23. **Page Replacement Algorithm:**

Page replacement algorithms are a crucial part of virtual memory management in operating systems. They determine which page to evict from physical memory (RAM) when there is a page fault (a requested page is not in RAM). Each algorithm has its characteristics, advantages, and trade-offs. Here, we'll explain the concept in detail:

- **Working Set:** The working set of a process is the set of pages it is actively using. Page replacement algorithms aim to keep the working set of each process in RAM to minimize page faults.

- **Page Table:** Each process maintains a page table that maps logical pages (used by the program) to physical frames (in RAM). The page table helps the OS and hardware manage memory.

- **Page Fault Handling:** When a page fault occurs, the OS must decide which page to replace. Various algorithms are used to make this choice:

- **FIFO (First-In, First-Out):** FIFO is a simple algorithm that replaces the oldest page in memory. While easy to implement, it may lead to the "Belady's Anomaly," where increasing the number of frames can result in more page faults.

- **LRU (Least Recently Used):** LRU replaces the page that hasn't been used for the longest time. It minimizes the number of page faults but can be computationally expensive to implement.

- **LFU (Least Frequently Used):** LFU replaces the page that has been used the least number of times. While it's effective in some cases, it can have difficulties handling changing access patterns.

- **Optimal (OPT):** OPT is an idealized algorithm that replaces the page that will not be used for the longest time in the future. While it provides the lowest possible number of page faults, it's not practical to implement since it requires knowledge of future page accesses.

- **Second-Chance (Clock):** Second-Chance combines aspects of FIFO and LRU. It examines the "age" of pages, considering whether they have been recently accessed. If not, they are candidates for replacement.

- **Page Replacement Policies:** The choice of a page replacement algorithm can significantly impact system performance. The optimal algorithm is challenging to implement, so most systems use approximations that balance between overhead and page fault rate.

- **Trade-Offs:** The main trade-off in page replacement is between the number of page faults (performance) and the overhead associated with tracking page accesses. No single algorithm is optimal for all scenarios, so the choice depends on system requirements and constraints.

Page replacement algorithms play a crucial role in memory management, ensuring efficient utilization of RAM and minimizing the impact of page faults on system performance.

24. **I/O Systems:**

Input/Output (I/O) systems are a critical part of operating systems, facilitating communication between the CPU, memory, and various peripheral devices. I/O systems enable data transfer to and from devices like disks, network interfaces, keyboards, and displays. Here, we'll explain I/O systems in detail:

- **Device Drivers:** Device drivers are software components that serve as intermediaries between the operating system and hardware devices. They provide a standardized interface for the OS to communicate with specific devices.

- **I/O Control:** I/O systems manage and control the flow of data between memory and devices. They ensure that data is correctly formatted, synchronized, and transferred to and from the correct locations.

- **Asynchronous I/O:** Asynchronous I/O allows programs to continue execution while waiting for I/O operations to complete. This can improve system responsiveness and efficiency.

- **Buffering:** Buffering involves temporarily storing data in memory before or after I/O operations. Buffers help smooth out differences in data transfer rates between devices and memory, reducing the impact of slow I/O operations.

- **Caching:** Caching is the practice of storing frequently accessed data in a cache to reduce the time it takes to retrieve the data. This is particularly important for speeding up disk I/O.

- **Interrupt Handling:** I/O systems rely on interrupts to notify the CPU when a device is ready for data transfer or when an error has occurred. Interrupt handlers are responsible for responding to these signals.

- **Synchronization:** I/O systems must coordinate the flow of data to and from devices and memory, ensuring that data is processed in the correct order and preventing conflicts.

- **File Systems:** I/O systems are closely tied to file systems, as they are responsible for reading and writing data to and from storage devices. File system drivers play a crucial role in managing I/O to files and directories.

- **Device Abstraction:** I/O systems provide device abstraction, allowing applications to interact with devices in a standardized way. This abstraction shields applications from hardware-specific details.
- **Network I/O:** In addition to local I/O, modern operating systems also support network I/O. This involves managing data transfer between network interfaces and network devices.
- **Performance Optimization:** I/O systems are optimized for performance, often using techniques like write-behind caching, read-ahead caching, and prefetching to reduce latency and increase throughput.
- **Error Handling:** Robust error handling is crucial in I/O systems to detect and recover from issues such as device failures, data corruption, and communication errors.

I/O systems are essential for the smooth operation of computer systems, ensuring that data can be reliably and efficiently moved between memory and peripheral devices. They play a vital role in overall system performance and user experience.

DECEMBER 2021

PART C — (3 × 10 = 30 marks)

Answer any THREE questions.

20. Discuss the essential properties of the types of Operating System.
21. Explain about System Call.
22. Point out the performance of demand paging in detail.
23. Discuss the various scheduling algorithms with examples.
24. Discuss about Directory Structure in detail.

Answers

20. **Essential Properties of Types of Operating Systems:**

Different types of operating systems exhibit distinct properties to cater to specific use cases and requirements. Here, we'll discuss the essential properties of a few key types of operating systems:

- **Real-Time Operating Systems (RTOS):** RTOS is designed for applications that demand predictable and timely responses. Key properties include determinism, where tasks

complete within specific time frames, and priority-based scheduling to meet deadlines. Safety-critical RTOSs emphasize fault tolerance.

- **Multi-User Operating Systems:** Multi-user operating systems provide concurrent access to multiple users. Essential properties include user isolation, security, and resource sharing. They employ authentication and access control mechanisms.

- **Multi-Tasking Operating Systems:** Multi-tasking OSes allow concurrent execution of multiple processes or tasks. Properties include context switching, efficient CPU scheduling, and memory protection to prevent processes from interfering with each other.

- **Distributed Operating Systems:** Distributed OSes focus on managing a network of interconnected computers as a single system. Key properties include transparency (hiding network complexities), fault tolerance (continuing operation despite failures), and efficient communication and resource sharing.

- **Mobile Operating Systems:** Mobile OSes are optimized for smartphones and tablets. Essential properties include touch screen support, power efficiency, app stores for easy application installation, and support for cellular communication and location services.

- **Embedded Operating Systems:** Embedded OSes run on devices with specific functions, like IoT devices and consumer electronics. Properties include minimal resource usage, real-time capabilities, and reliability. They often run on specialized hardware.

- **Network Operating Systems:** Network OSes manage network-related tasks, such as routing, data sharing, and network protocol handling. Key properties include network device support, protocol implementation, and remote administration.

- **Single-User Operating Systems:** These are simple OSes for legacy or basic systems. Essential properties include providing a simple user interface, minimal resource requirements, and support for basic I/O operations.

The essential properties of an operating system determine its suitability for specific use cases and applications. Choosing the right type of OS is crucial for ensuring optimal system performance and functionality.

21. **System Call:**

A system call (syscall) is a fundamental interface between application programs and the operating system. System calls allow user-level processes to request services or functionality from the underlying kernel, which runs in privileged mode. Here, we'll discuss system calls in detail:

- **Interface:** System calls provide a well-defined interface between user-level applications and the operating system. This interface typically includes functions, system call numbers, and parameters for requesting specific services.

- **Invocation:** To make a system call, a program typically uses programming language constructs or library functions. For example, in C or C++, functions like `open()`, `read()`, and `write()` are used for file operations, and these functions ultimately invoke system calls.
- **Privilege Transition:** When a program makes a system call, it triggers a privilege transition from user mode to kernel mode. This transition allows the program to access protected resources and execute privileged operations.
- **Functionality:** System calls cover a wide range of functionality, including file operations, process management, memory allocation, I/O operations, networking, and more. Common examples include `fork()` for process creation, `exec()` for executing a new program, and `socket()` for network communication.
- **Error Handling:** System calls can return error codes to indicate exceptional conditions, such as file not found or insufficient permissions. Application programs must check these codes to handle errors appropriately.
- **Security:** System calls play a crucial role in maintaining security and access control. They ensure that user-level processes cannot directly access privileged resources or perform unauthorized actions.
- **User and Kernel Space:** System calls bridge the gap between user space, where application code runs, and kernel space, where the operating system operates. This separation enforces security and isolation.

System calls are a core component of the interaction between software and the operating system. They allow programs to harness the power of the underlying hardware and OS services while maintaining a secure and controlled environment.

22. **Performance of Demand Paging:**

Demand paging is a memory management technique used in virtual memory systems to optimize RAM usage. It has several performance advantages and some considerations:

- **Reduced Memory Requirements:** Demand paging allows the operating system to load only the necessary pages of a process into RAM. This reduces memory requirements, enabling more processes to run concurrently.
- **Efficient Use of RAM:** In a demand paging system, only actively used pages are resident in RAM. This means that processes are less likely to suffer from thrashing, where constant page swapping hampers performance.
- **Isolation:** Demand paging provides memory isolation between processes. Each process has its virtual address space, making it less susceptible to interference from other processes.
- **Memory Protection:** Demand paging enforces memory protection, preventing one process from accessing the memory of another process. Unauthorized access is detected and leads to a segmentation fault or similar error.

- **Copy-on-Write:** Many demand paging systems implement copy-on-write, which allows multiple processes to share the same physical memory until one of them modifies it. This reduces memory usage and page copying overhead.

- **Adaptive Page Replacement:** Demand paging systems use page replacement algorithms (e.g., LRU, FIFO) to choose which pages to evict when new pages are demanded. Adaptive algorithms improve the system's ability to retain frequently accessed pages.

- **Effective Use of Secondary Storage:** Demand paging effectively uses secondary storage (e.g., disk) as an extension of RAM. It ensures that rarely used pages are swapped out, freeing up RAM for more active pages.

Despite these advantages, demand paging can introduce performance overhead, especially when page faults occur. Excessive page faults can lead to disk I/O bottlenecks, which impact system responsiveness. To mitigate this, modern operating systems employ various page replacement algorithms and memory management strategies to balance the trade-off between RAM usage and page fault frequency.

Overall, demand paging is a valuable memory management technique that significantly improves memory utilization and enables efficient multitasking. Its performance depends on factors like the page replacement algorithm, the size of RAM, and the I/O performance of secondary storage.

23. **Scheduling Algorithms with Examples:**

CPU scheduling is a critical aspect of operating systems, responsible for determining the order in which processes are executed on the CPU. Various scheduling algorithms have been developed, each with its characteristics. Here, we'll discuss some scheduling algorithms with examples:

- **First-Come, First-Served (FCFS):** FCFS is a non-preemptive scheduling algorithm where processes are executed in the order they arrive. It's simple but can lead to the "convoy effect," where a long process holds up shorter ones. Example: Consider three processes arriving in the order P1, P2, P3. They are executed in the same order.

- **Shortest Job First (SJF):** SJF is a non-preemptive or preemptive scheduling algorithm where the process with the shortest burst time is selected next. It minimizes average waiting time and turnaround time but requires knowledge of process execution times. Example: Suppose processes P1, P2, and P3 have burst times of 5, 2, and 8 milliseconds. They are executed in the order P2, P1, P3, as P2 has the shortest burst time.

- **Round Robin (RR):** RR is a preemptive scheduling algorithm where each process is assigned a fixed time quantum. When the time quantum expires, the process is moved to the back of the queue, allowing the next process to execute. It's fair but can lead to overhead due to context switches. Example: Processes P1, P2, P3 are scheduled with a time quantum of 4 milliseconds. They execute in a circular fashion: P1, P2, P3, P1, P2, P3, and so on.

- **Priority Scheduling:** Priority scheduling assigns a priority to each process, and the process with the highest priority is executed first. It can be either preemptive or non-preemptive. Care is needed to prevent priority inversion issues. Example: Suppose processes P1, P2, and P3 have priorities 3, 2, and 1, respectively. They are executed in order of priority.

- **Multilevel Queue Scheduling:** Multilevel queue scheduling divides processes into multiple queues with different priorities. Each queue may use a different scheduling algorithm. For example, interactive processes may be in a high-priority queue, while batch processes are in a low-priority queue.

- **Multilevel Feedback Queue Scheduling:** Similar to multilevel queue scheduling, but with the ability to move processes between queues based on their behavior. This allows for dynamic adjustment of priorities.

- **Priority Inversion:** Priority inversion can occur when a higher-priority process is indirectly delayed by a lower-priority process. It's often solved using techniques like priority inheritance.

Scheduling algorithms are chosen based on the system's requirements and goals, balancing factors like fairness, efficiency, and response time. The choice of algorithm can significantly impact system performance and user experience.

24. **Directory Structure:**

Directory structures are an essential part of file systems that organize files and directories on storage devices. There are several directory structure designs, each with its characteristics. Here, we'll discuss directory structures in detail:

- **Single-Level Directory:** In a single-level directory structure, all files are stored in a single directory. This structure is simple but lacks organization, making it challenging to manage large numbers of files. It's commonly found in older operating systems.

- **Two-Level Directory:** In a two-level directory structure, files are organized into user directories within the root directory. Each user has their directory, providing better organization and isolation. However, there can be naming conflicts between users. Example: User directories could be /root/user1, /root/user2, etc.

- **Tree-Structured Directory:** A tree-structured directory system allows for a hierarchical organization of files and directories. Each directory can contain both files and subdirectories. This structure is widely used in modern file systems and provides excellent organization and isolation. Example: /root/ |--- Documents | |--- File1.txt | |--- File2.txt |--- Pictures |
|--- Photo1.jpg | |--- Photo2.jpg

- **Acyclic-Graph Directory:** In this structure, a directory can have multiple parent directories, forming a directed acyclic graph (DAG). This design allows for flexible organization but requires more complex management to avoid loops in the graph.

- **General Graph Directory:** In a general graph directory structure, directories can have multiple parents, and cycles are allowed in the directory graph. This is a more complex and less common structure due to the potential for confusion and management challenges.

- **File Allocation Table (FAT):** The File Allocation Table structure uses a table to keep track of file locations on the disk. It's commonly used in the FAT file system family, found in older Windows versions and some removable storage devices.

- **Indexed Allocation:** In indexed allocation, each file has an associated index block that contains pointers to data blocks. This allows for efficient direct access to data but may require multiple reads to retrieve file data.

- **Clustered or Contiguous Directory:** In a clustered or contiguous directory structure, files are stored in contiguous clusters on the disk. This design minimizes disk head movement for sequential access but can lead to fragmentation.

Directory structures play a crucial role in file system organization, making it easier for users and applications to locate and manage files. The choice of structure depends on the file system's requirements and the intended use cases.